

## **Today:**

- Modifications to indexed-sequential files.
- Secondary indexes, inverted indexes.
- B-trees: a clean way to manage multilevel indexes.

## **Soon:**

- Hashing, another powerful indexing technique.
- Multidimensional index structures.

## DB Modifications

When we insert or delete on the data file, here are the primitive actions we might take:

1. Create or destroy an empty block in the sequence of blocks belonging to the sequential file.
2. Create or destroy an overflow block.
3. Insert a record into a block that has room.
4. Delete a record.
5. Slide a record to an adjacent block.


## Effect of Primitive Actions on Index File

Action	Dense	Sparse
Create/destroy empty overflow block	none	none
Create empty seq. block	none	insert
Destroy empty seq. block	none	delete
Insert record	insert	update(?)
Delete record	delete	update(?)
Slide record	update	update(?)


## Options

- Compact data and/or index blocks (all empty space at end).
  - ❖ Only need to record beginning of available space, rather than empty/full for each record-slot.
  - ❖ Compaction a problem if pointers from outside index.
- Add sequential blocks or overflow blocks when space needed.
  - ❖ Overflow blocks require no change to index, but make sparse indexes “sparser.”
- Redistribute records locally or always go for a new block.
  - ❖ Redistribution is extra effort, but keeps blocks more full.


## Example: Delete 30 With Dense Index




## Example: Delete 30 With Sparse Index



## Example: Insert 15 With Sparse Index — Redistribute




## Use Overflow Block Instead





## Secondary Indexes

- SKS says *primary index* is an index on a sorted file.
- I prefer to consider any index that “controls” the placement of records to be primary, e.g., hash table.
- *Secondary index* = index that does not control placement, surely not on a file sorted by its search key.
  - ❖ Sparse, secondary index makes no sense.
  - ❖ Usually, search key is not a “key.”



## Indirect Buckets


To avoid repeating keys in index, use a level of indirection, called *buckets*.

- Additional advantage: allows intersection of sets of records without looking at records themselves.

## Example

Movies(title, year, length, studioName);  
secondary indexes on studioName and year.

```
SELECT title
FROM Movies
WHERE studioName = 'Disney' AND
       year = 1995;
```



Buckets  
for  
studio

Movie tuples

Buckets  
for  
year




Disney	

Studio  
index

1995	

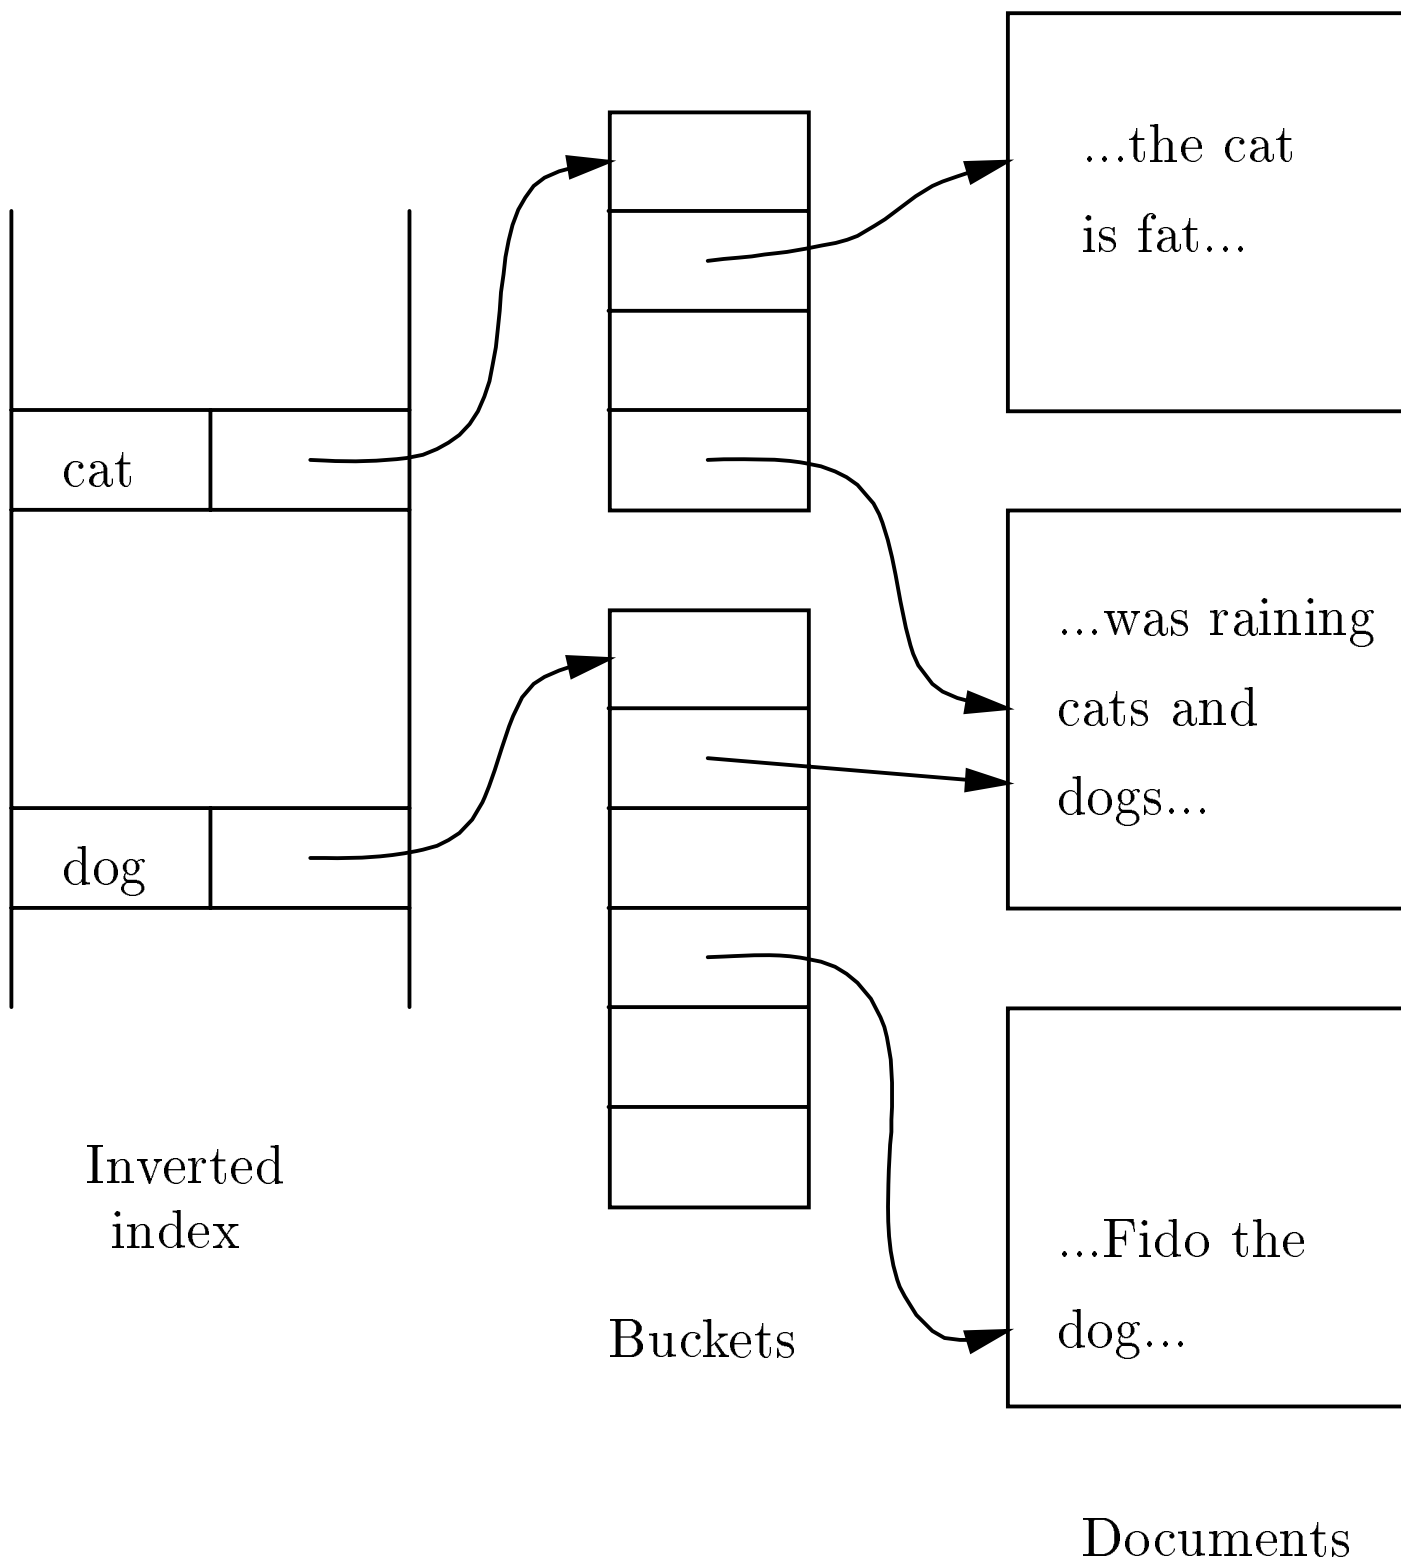
Year  
index

## **Inverted Indexes**

Similar (to secondary indexes) idea from information-retrieval community, but:

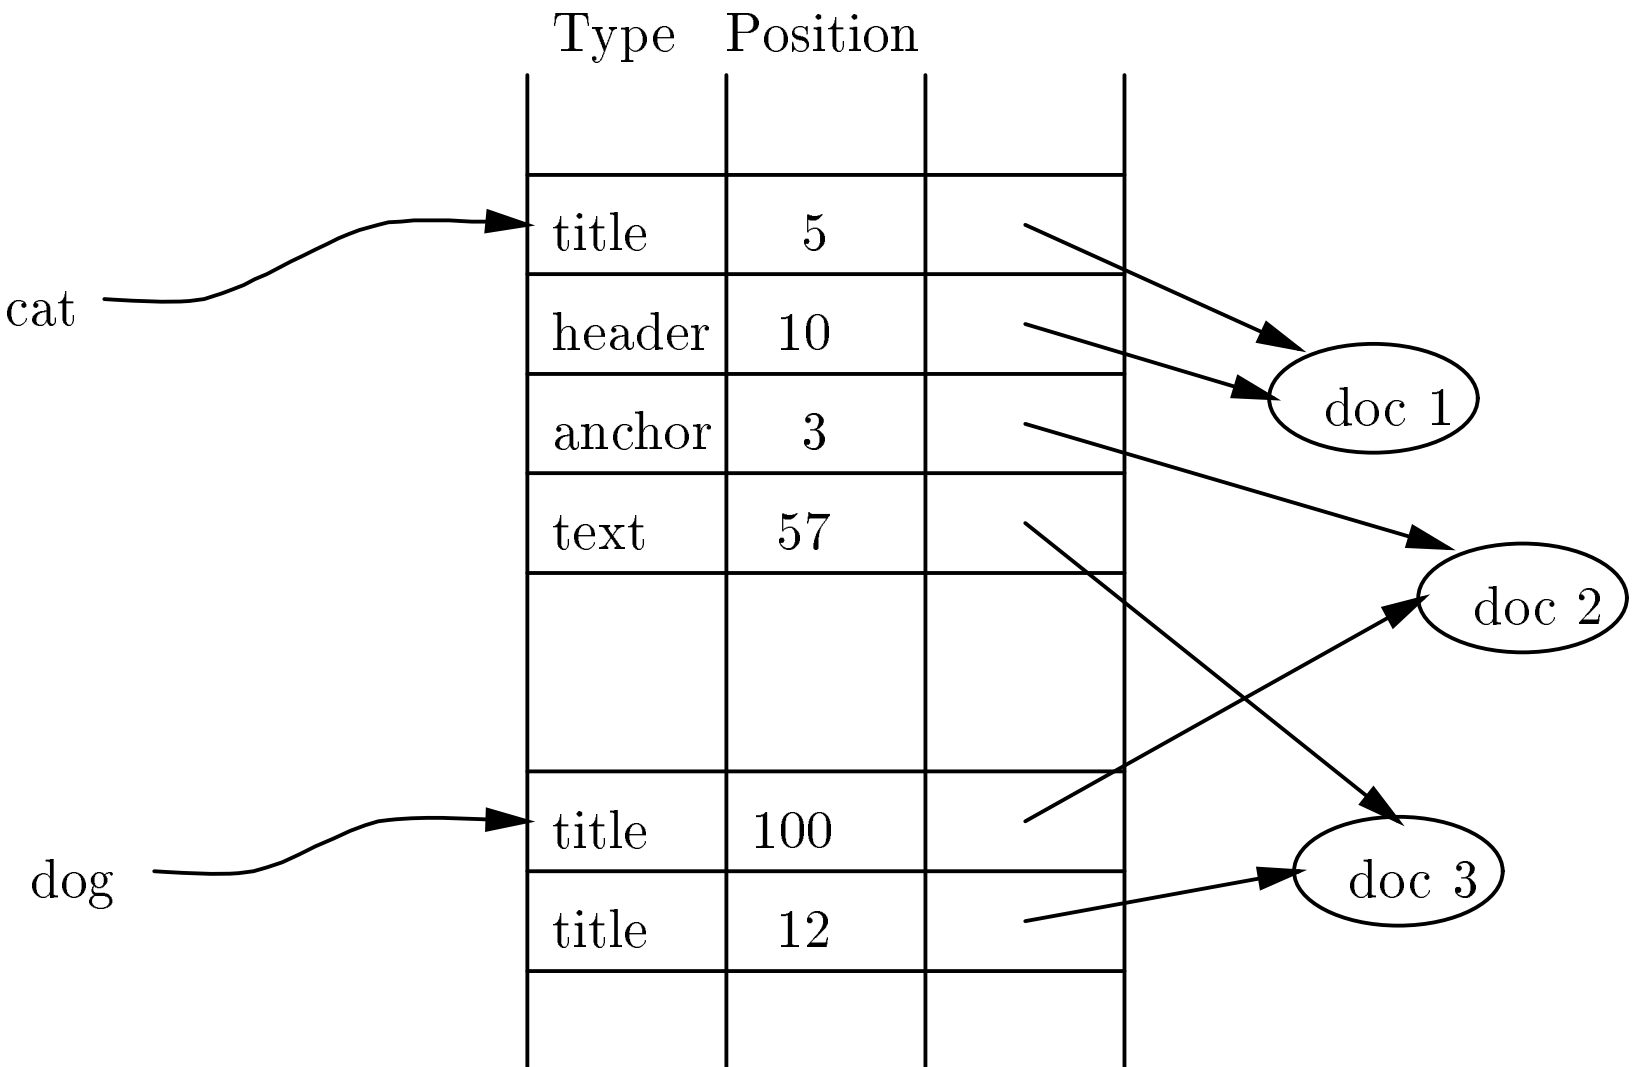
- Record  $\rightarrow$  document.
- Search-key value of record  $\rightarrow$  presence of a word in a document.

Usually used with “buckets.”



## Additional Information in Buckets

Can extend bucket to include role, position of word, e.g.






## B-Trees

Generalizes multilevel index.

- Number of levels varies with size of data file, but is often 3.
- *B+ tree* = form we'll discuss.
  - ❖ All nodes have same format:  $n$  keys,  $n + 1$  pointers.
- Useful for primary, secondary indexes, primary keys, nonkeys.


## Leaves

- One pointer to next leaf.
- $n$  key-pointer pairs for records of data file.
- At least half of these (round up) occupied.



## Interior Nodes

- $n$  keys form the divisions among  $n+1$  subtrees.
  - ❖ Key  $i$  is least key reachable from  $(i + 1)$ st child.
- At least  $n/2$  (round down) keys used, and one more pointer than key is used.
  - ❖ Exception: root may have only 2 children, one key.



## If There are Duplicate Keys

Key  $i$  is least *new* key reachable from  $(i+1)$ st child.

- Exception: the sole key if there is only one key in that entire subtree.

## Lookup in B+ Tree

- Start at root.
- Until you reach a leaf, follow the pointer that could lead to the key you want.
- Search that leaf (and leaves to the right if duplicates are possible).

## B+ Tree Insertion

- Search for the key being inserted.
- If there is room for another key-pointer at that leaf, insert there.
- If no room, split leaf.
  - ❖ Split of leaf looks like insert of child at level above.
  - ❖ Thus, recursive splitting all the way up the tree is possible.
  - ❖ Be careful to adjust keys as tree changes.

## B+ Tree Deletion

- Search for key being deleted.
- If found, delete.
- If the lower limit on occupancy is violated:
  - ❖ First look for an adjacent leaf that is above lower limit; “steal” a key-pointer pair from that leaf.
  - ❖ If none, then there must be two adjacent leaves, one at minimum, one below minimum. Just enough to merge nodes.
  - ❖ Merger looks like delete above, so recursive deletion possible.
  - ❖ Again, make sure keys are adjusted above.
- Sometimes, it is OK to allow a B-tree leaf to become subminimum — no mergers.